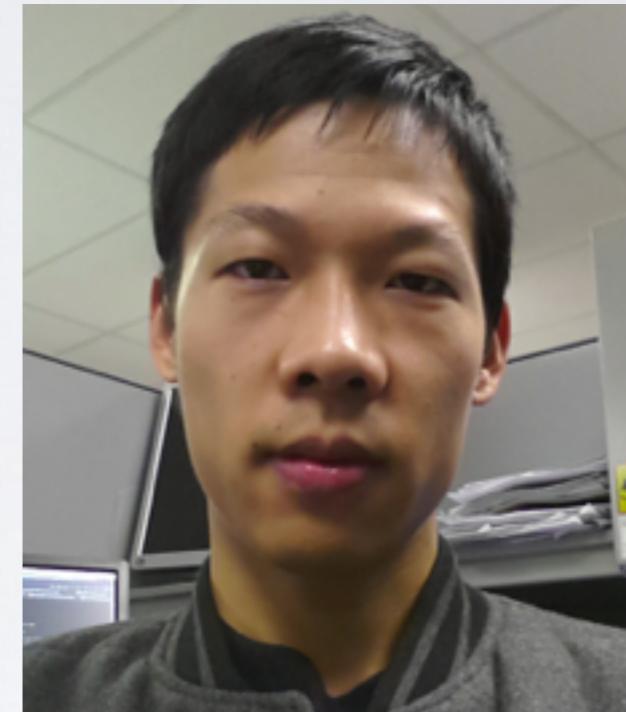


# FLINK: UNIFIED STREAM ENGINE

Li Chengxiang

# ABOUT ME

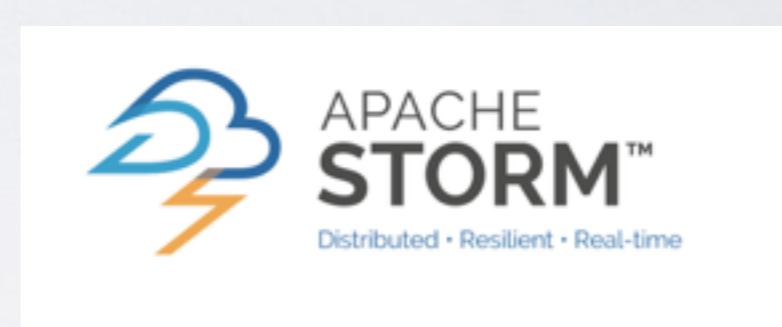
- Apache Hive Committer
- Apache Flink Committer
- Focus on large scale distributed computer engine and distributed SQL engine.



# AGENDA

- Batch and Stream
- Core feature of Stream
- Advanced Flink Stream feature

# BATCH AND STREAM



# IT'S ALL ABOUT DATA

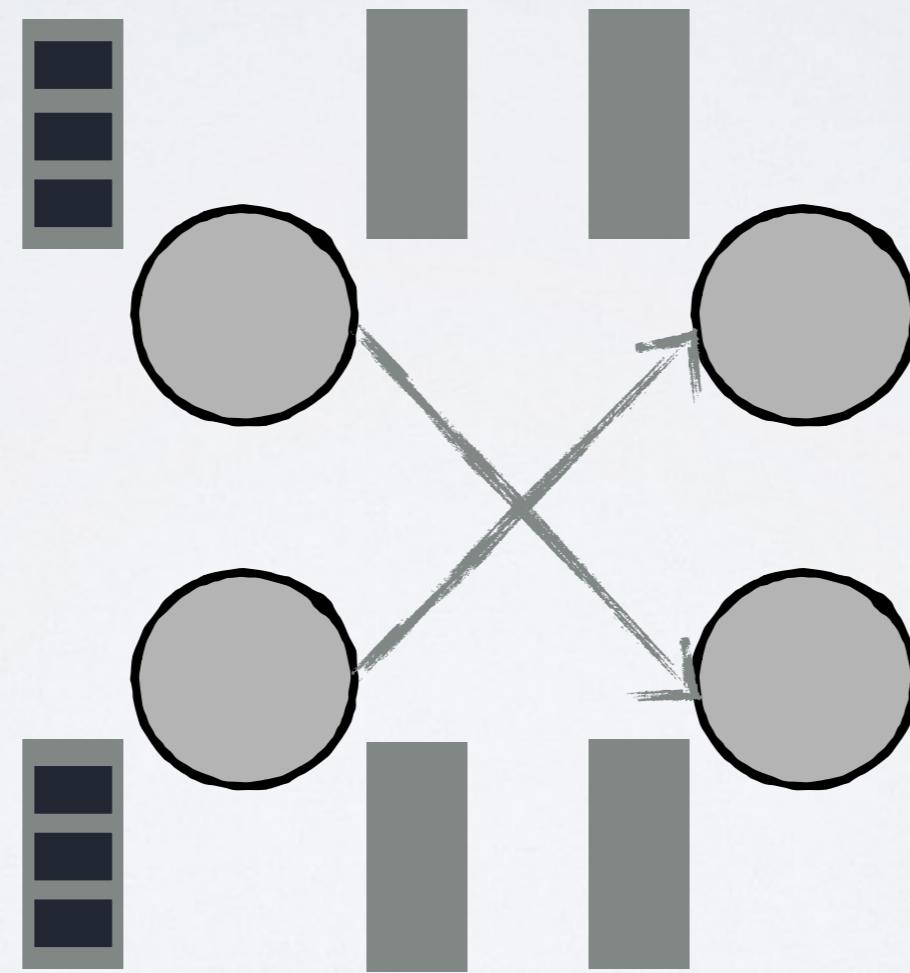


Bounded Data

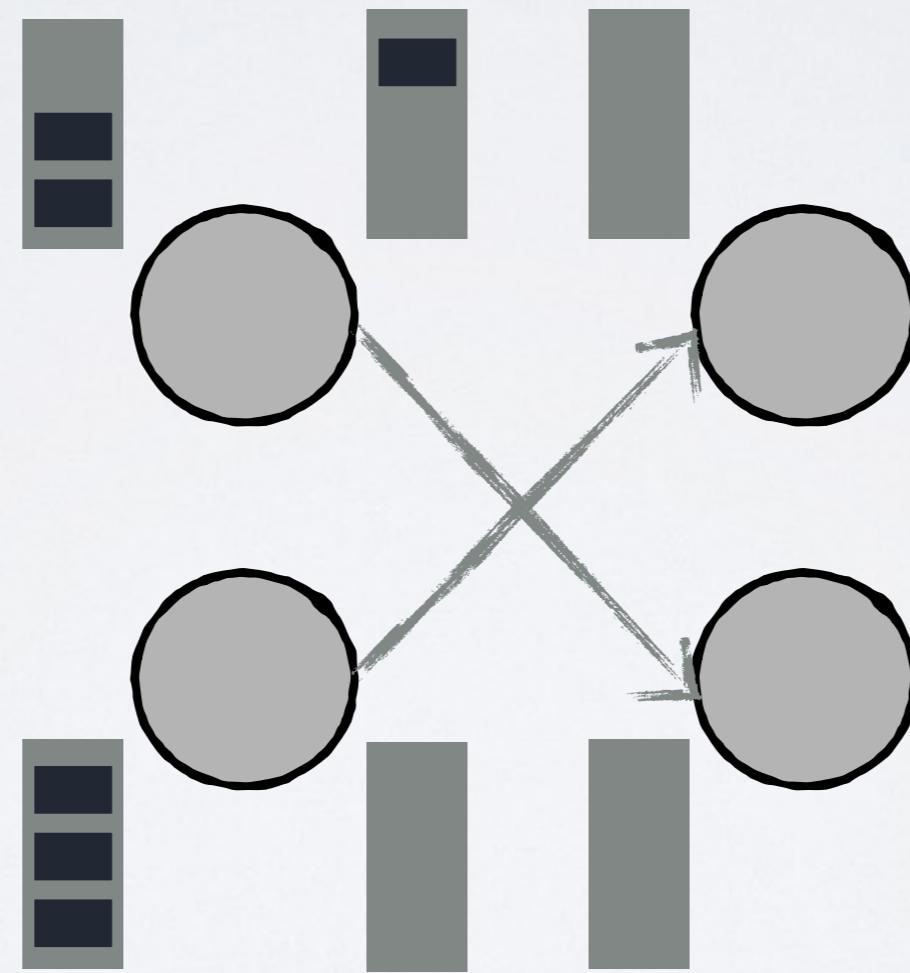


Unbounded Data

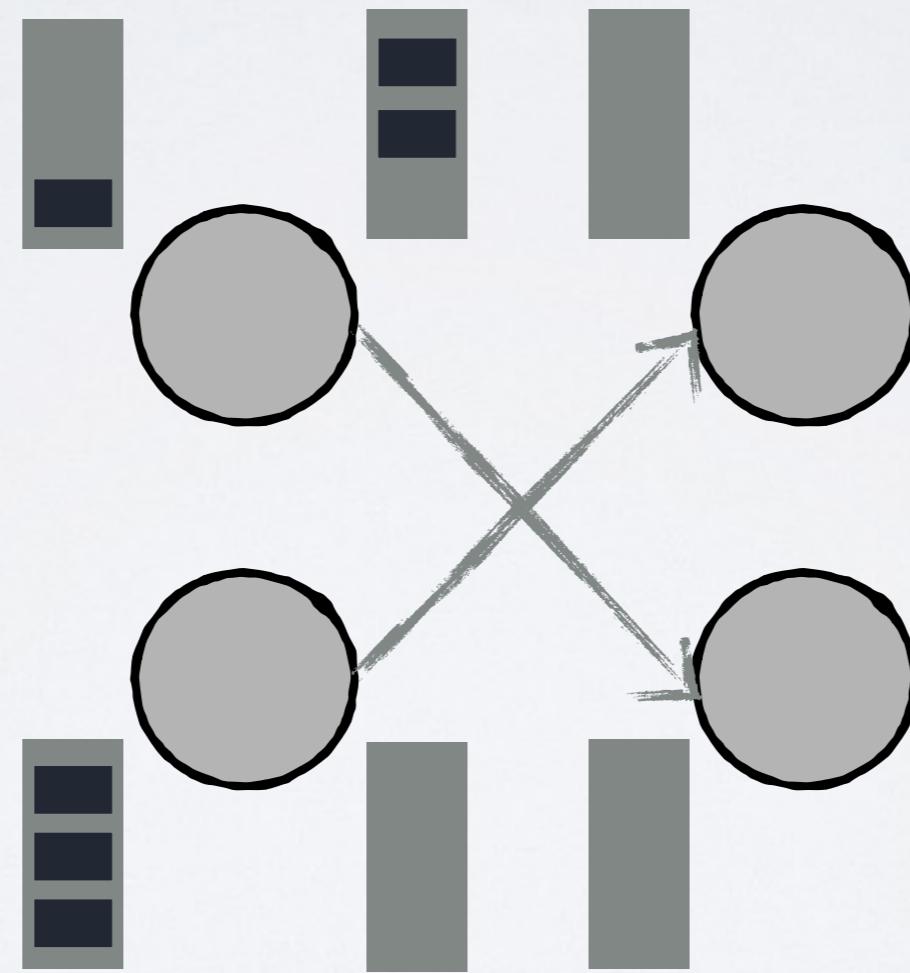
# BATCH ENGINE



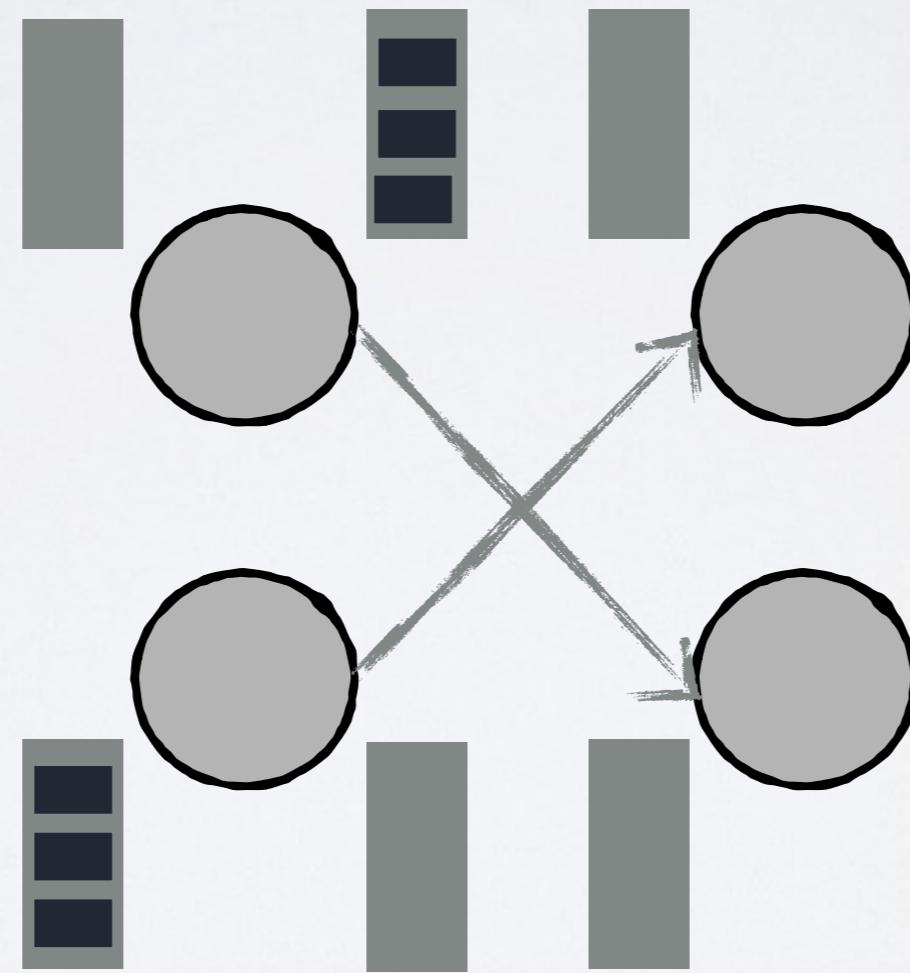
# BATCH ENGINE



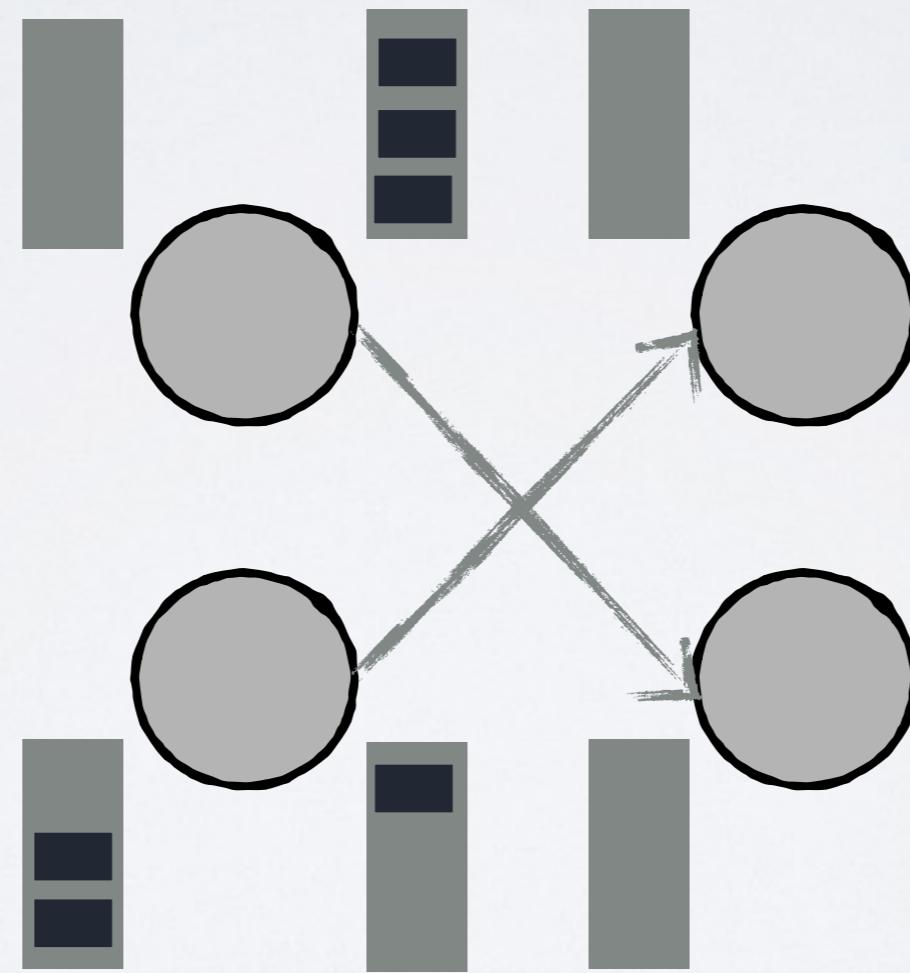
# BATCH ENGINE



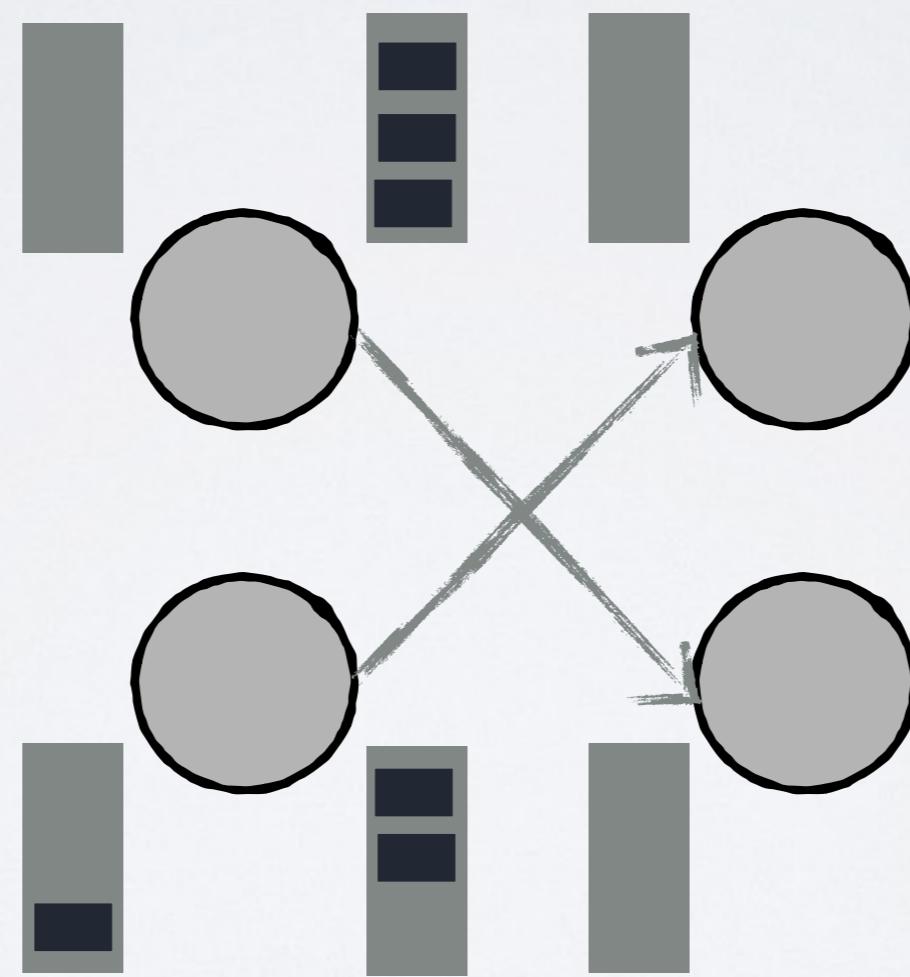
# BATCH ENGINE



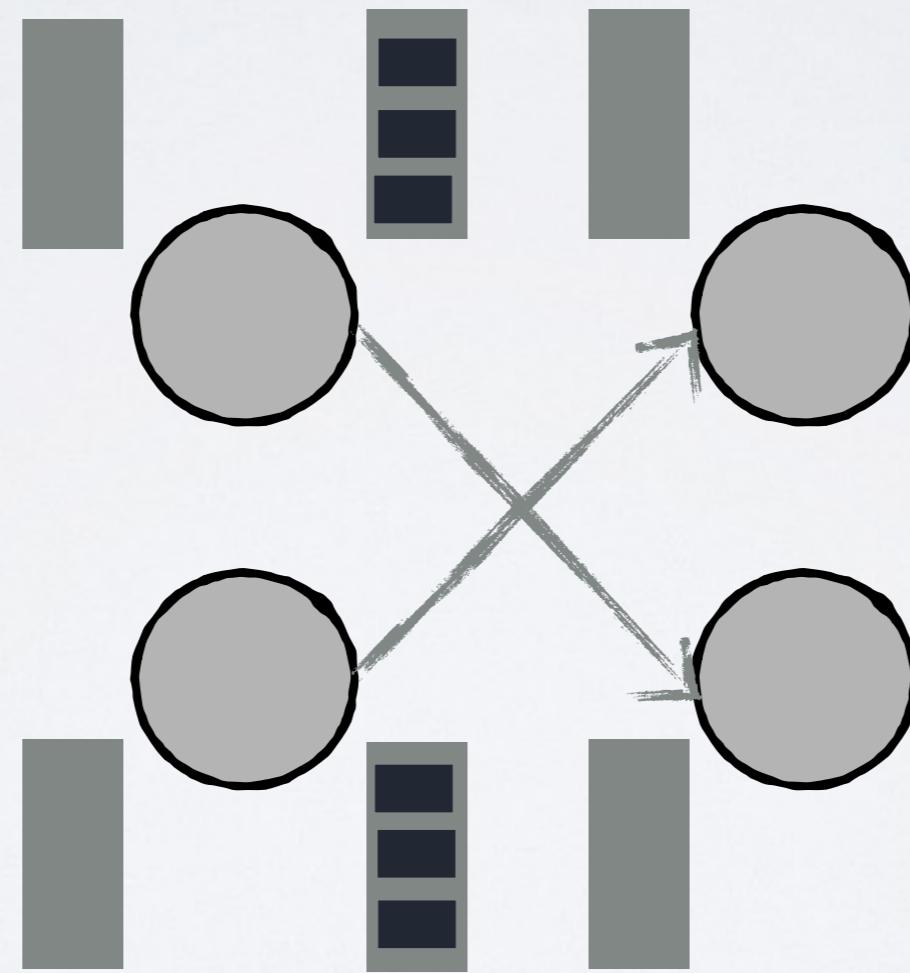
# BATCH ENGINE



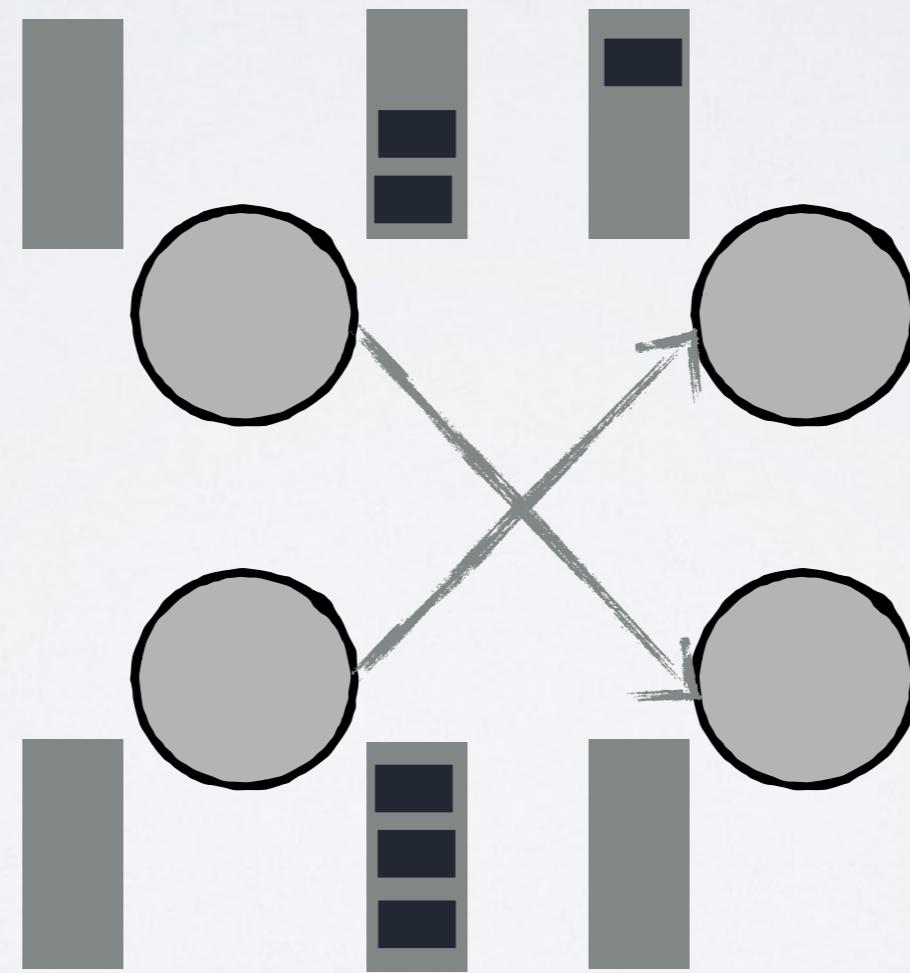
# BATCH ENGINE



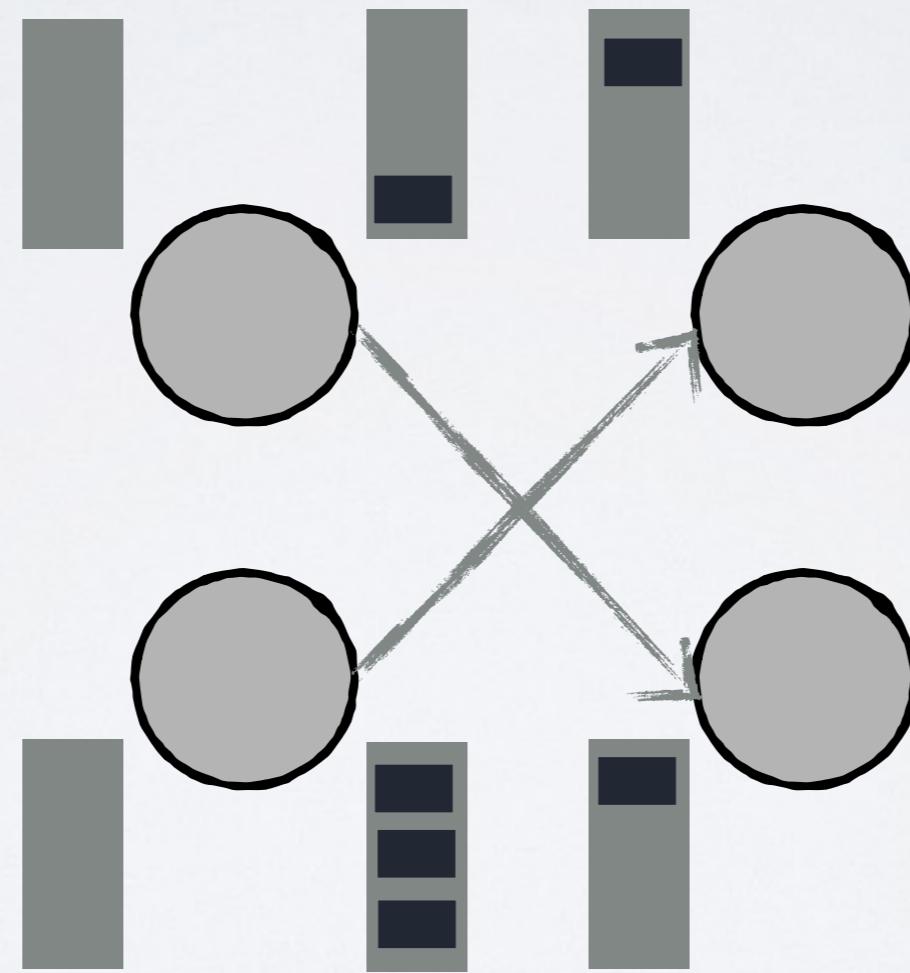
# BATCH ENGINE



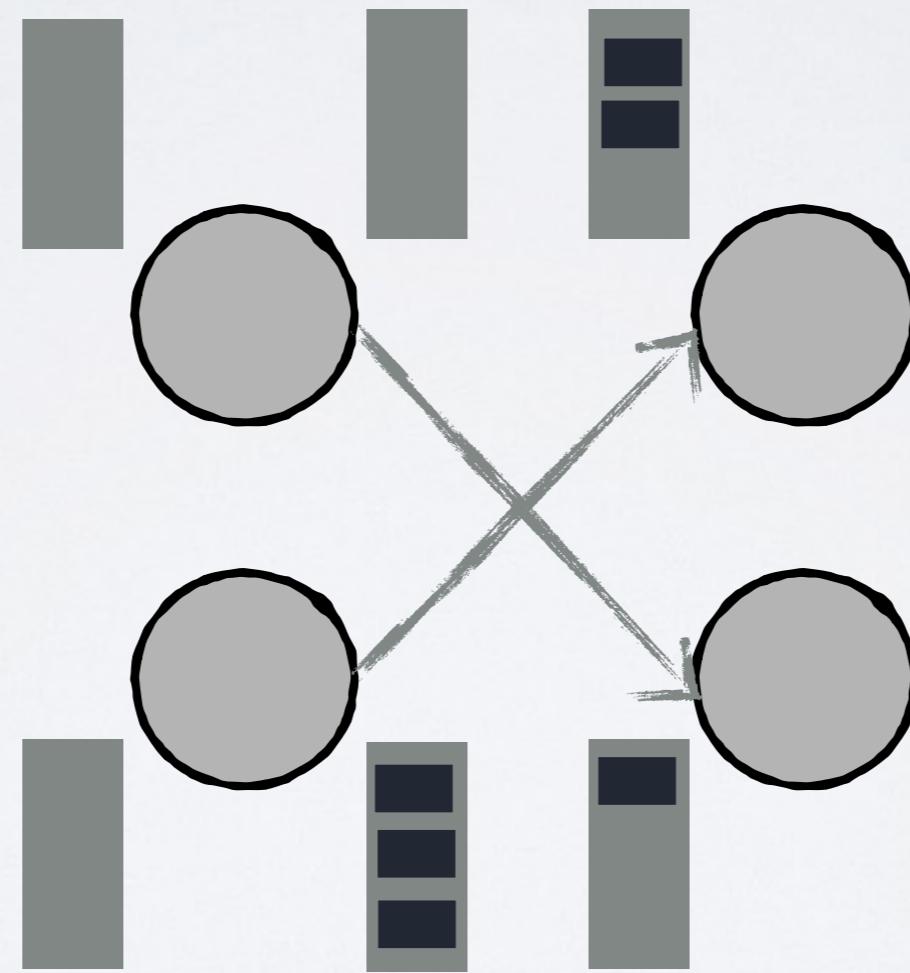
# BATCH ENGINE



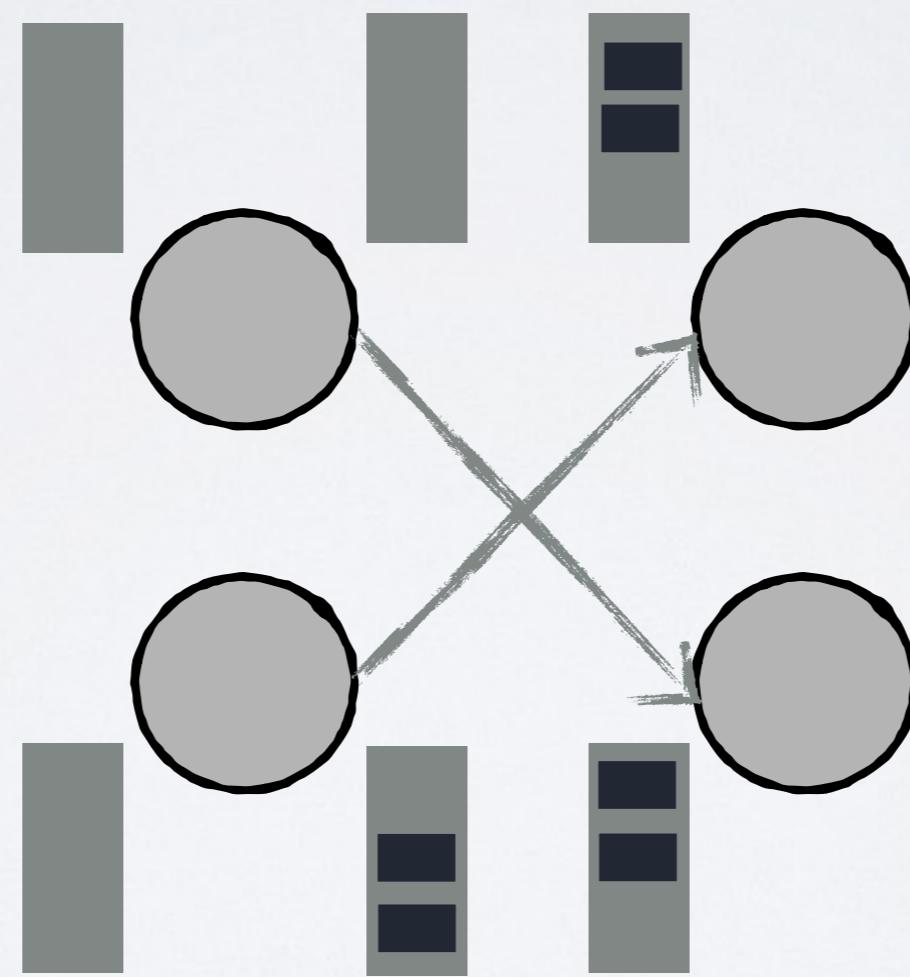
# BATCH ENGINE



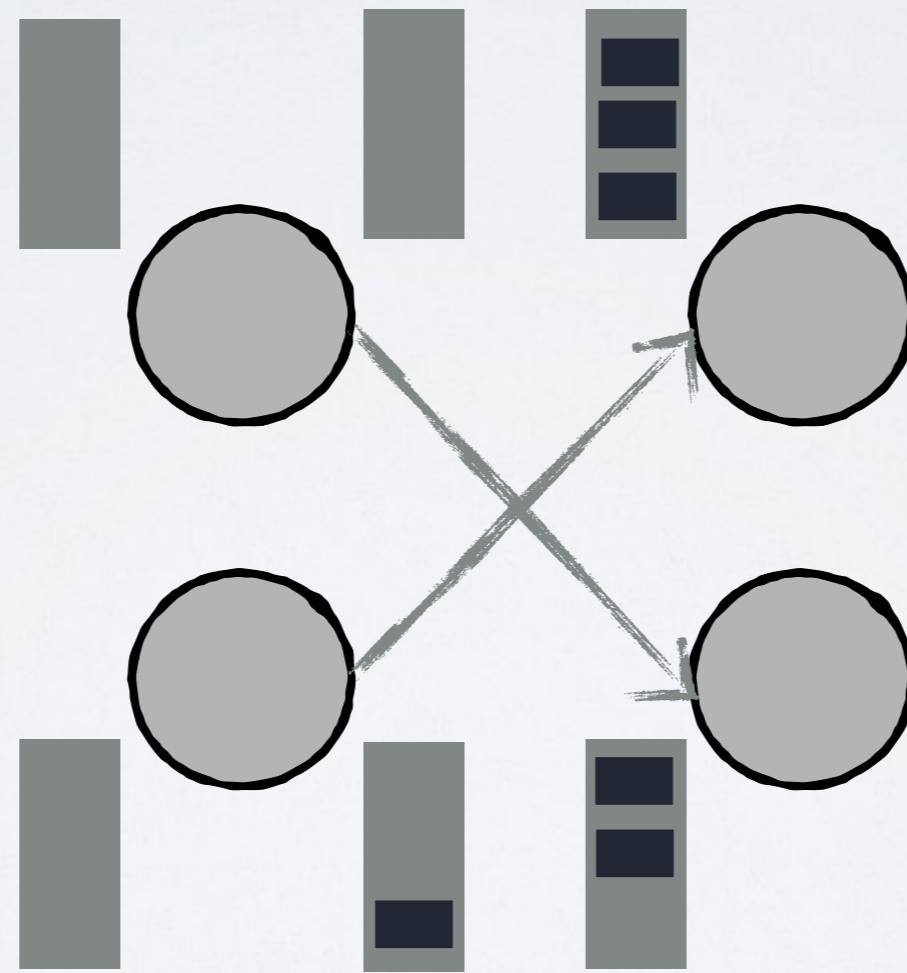
# BATCH ENGINE



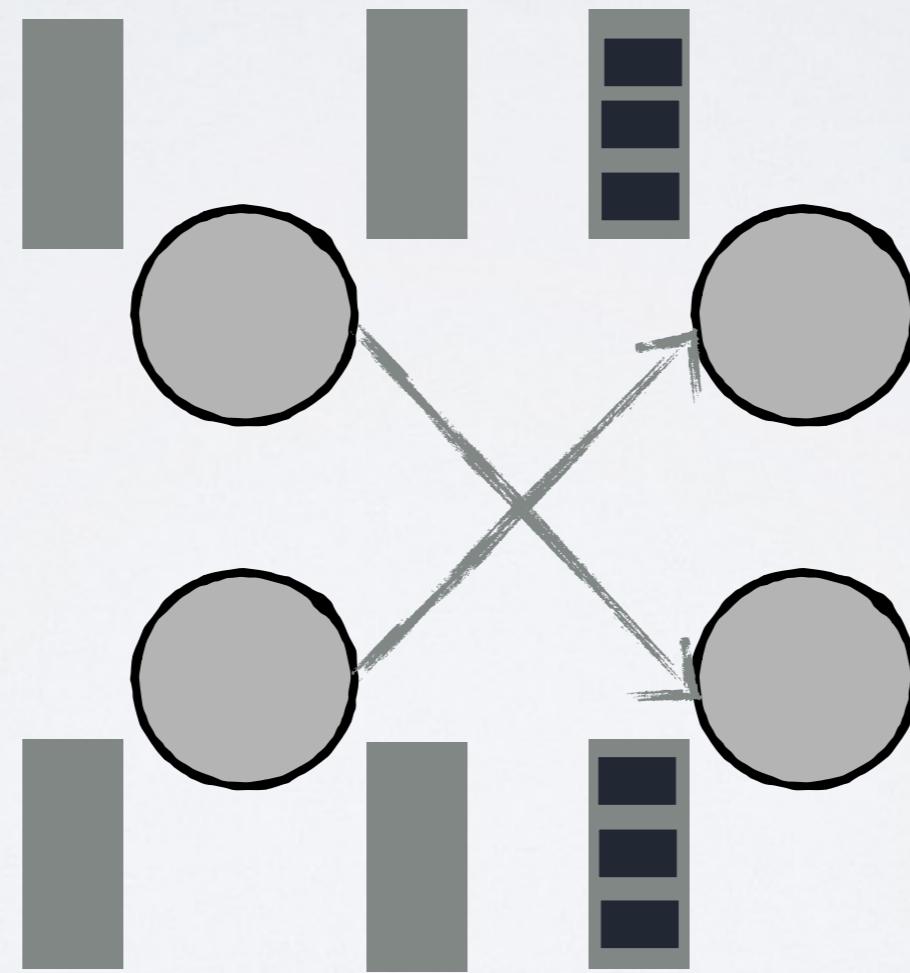
# BATCH ENGINE



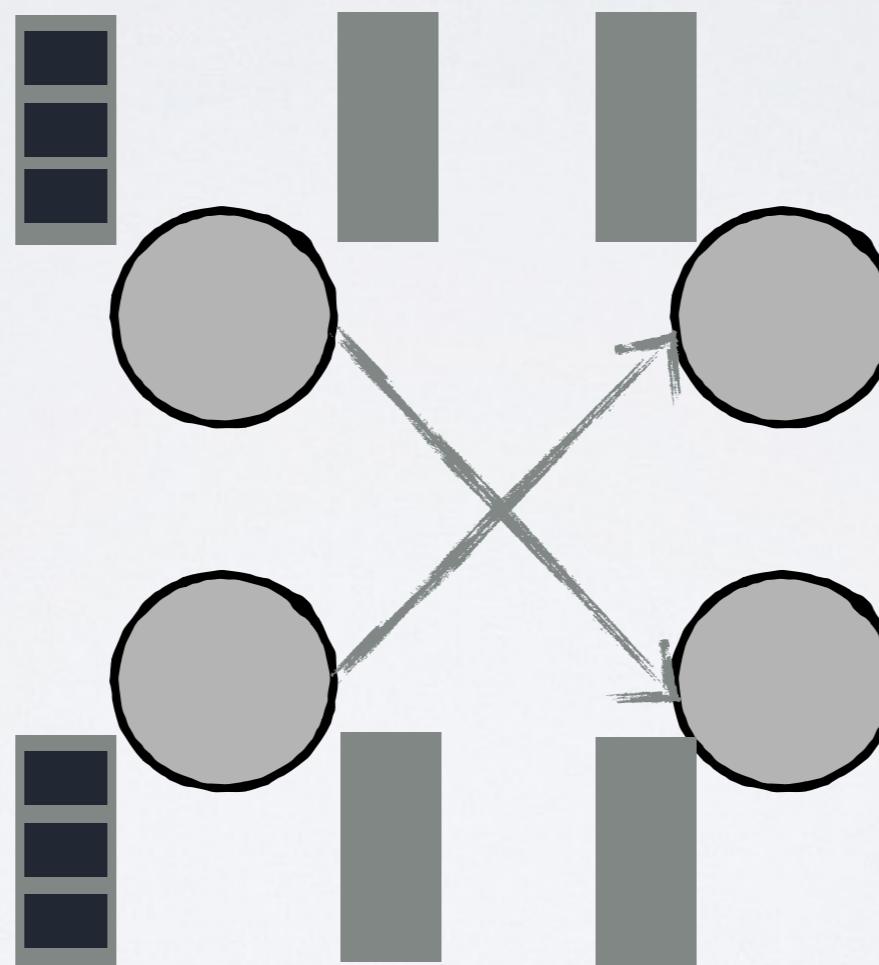
# BATCH ENGINE



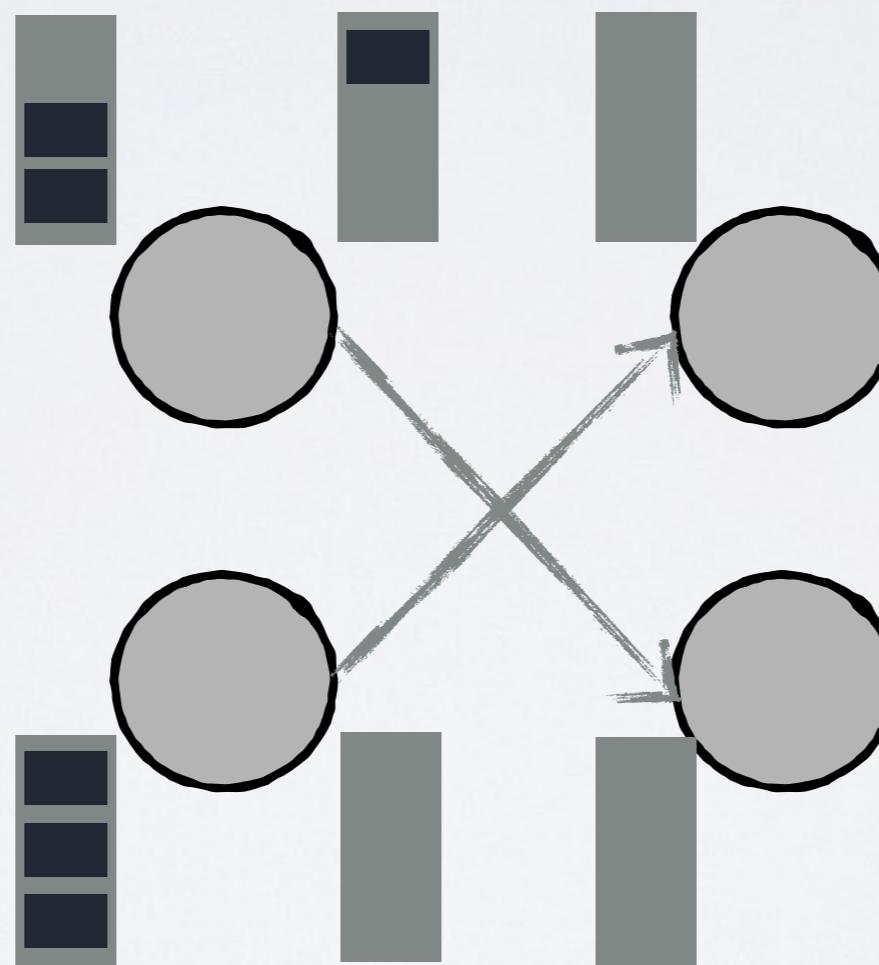
# BATCH ENGINE



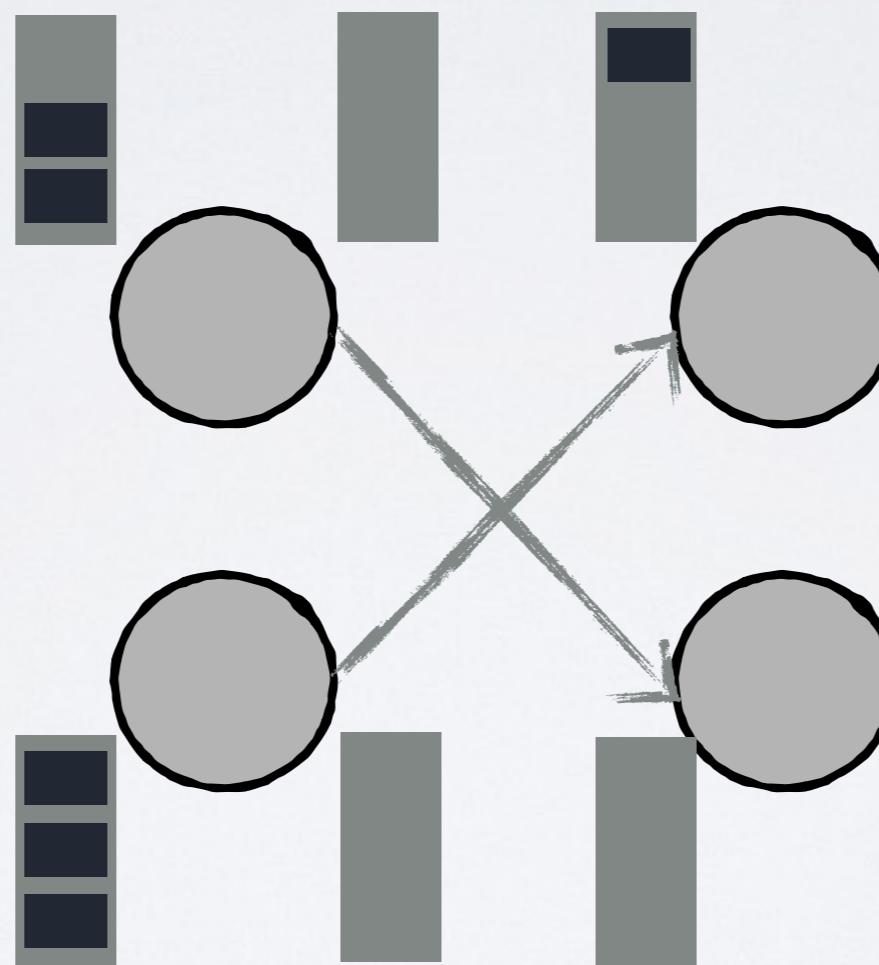
# STREAM ENGINE



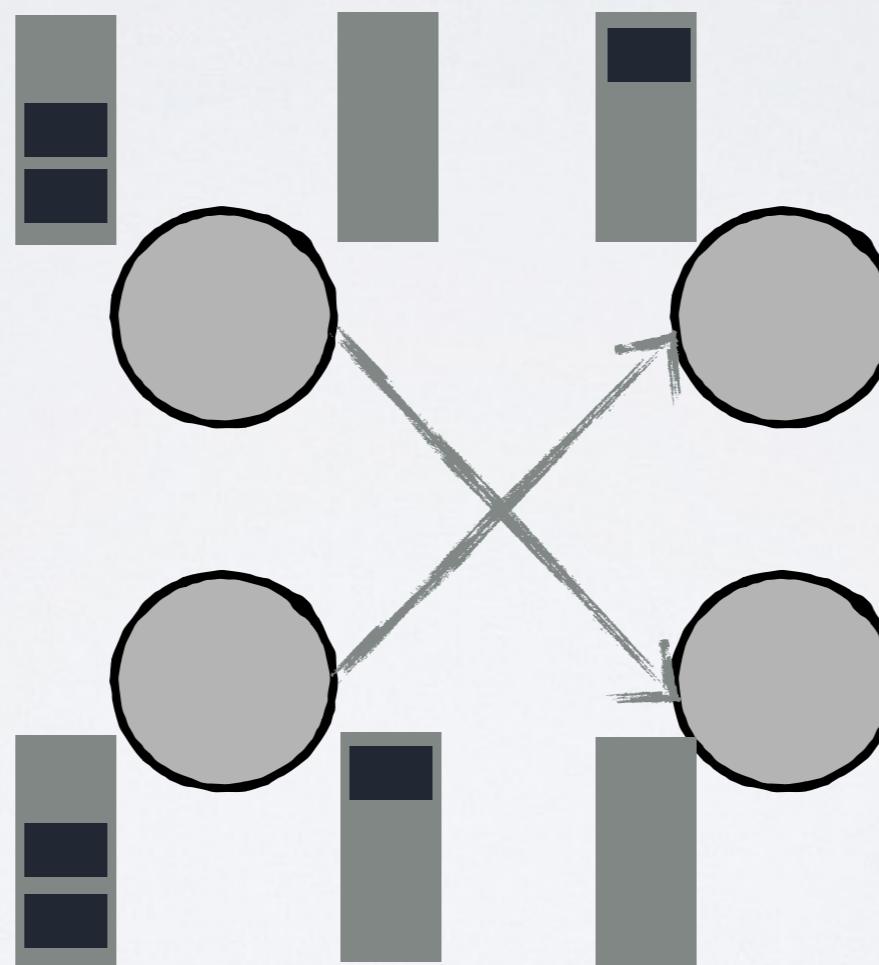
# STREAM ENGINE



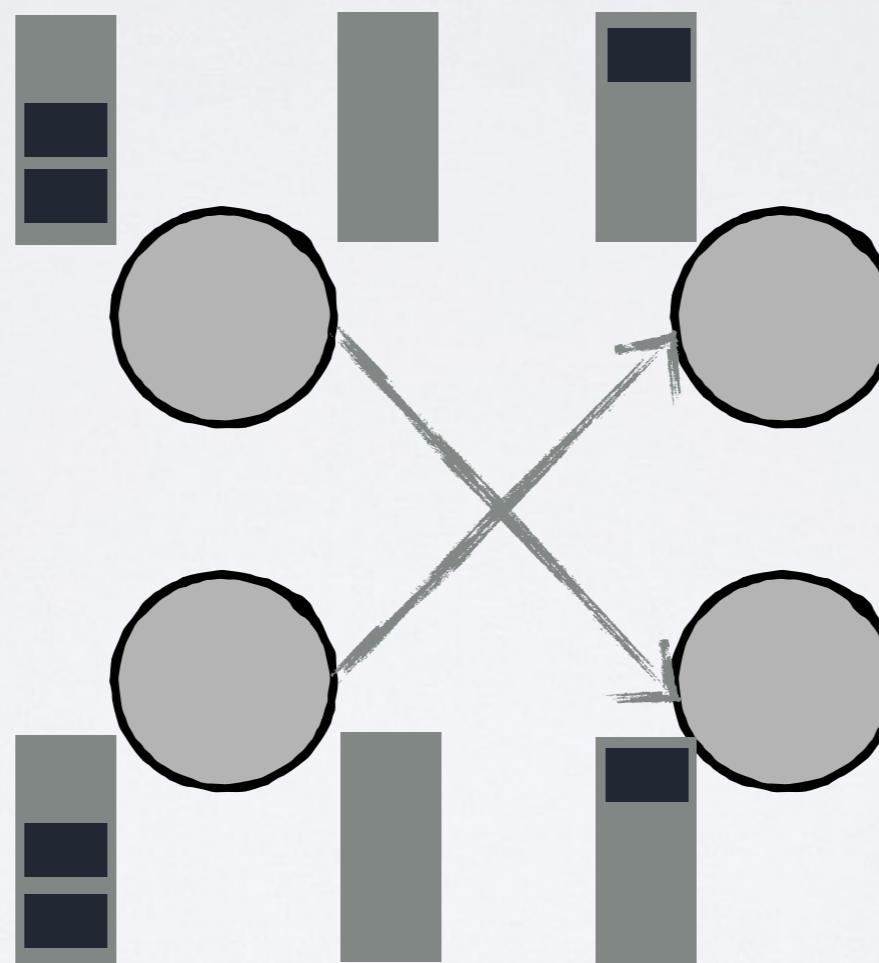
# STREAM ENGINE



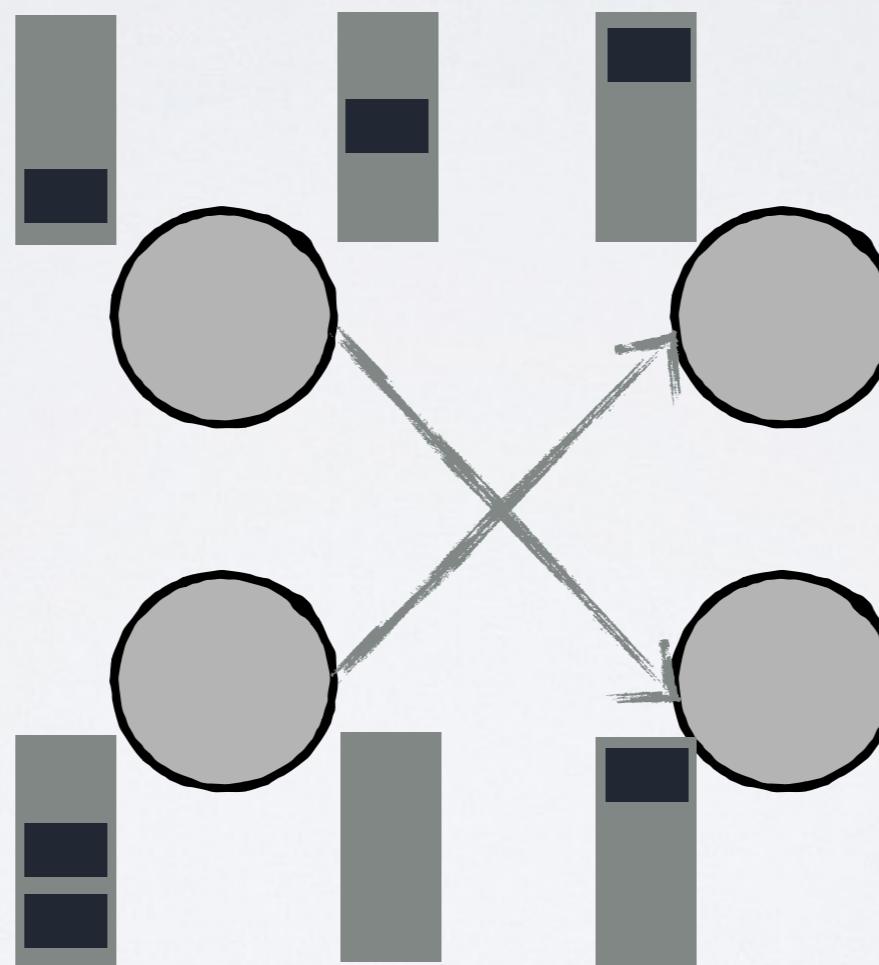
# STREAM ENGINE



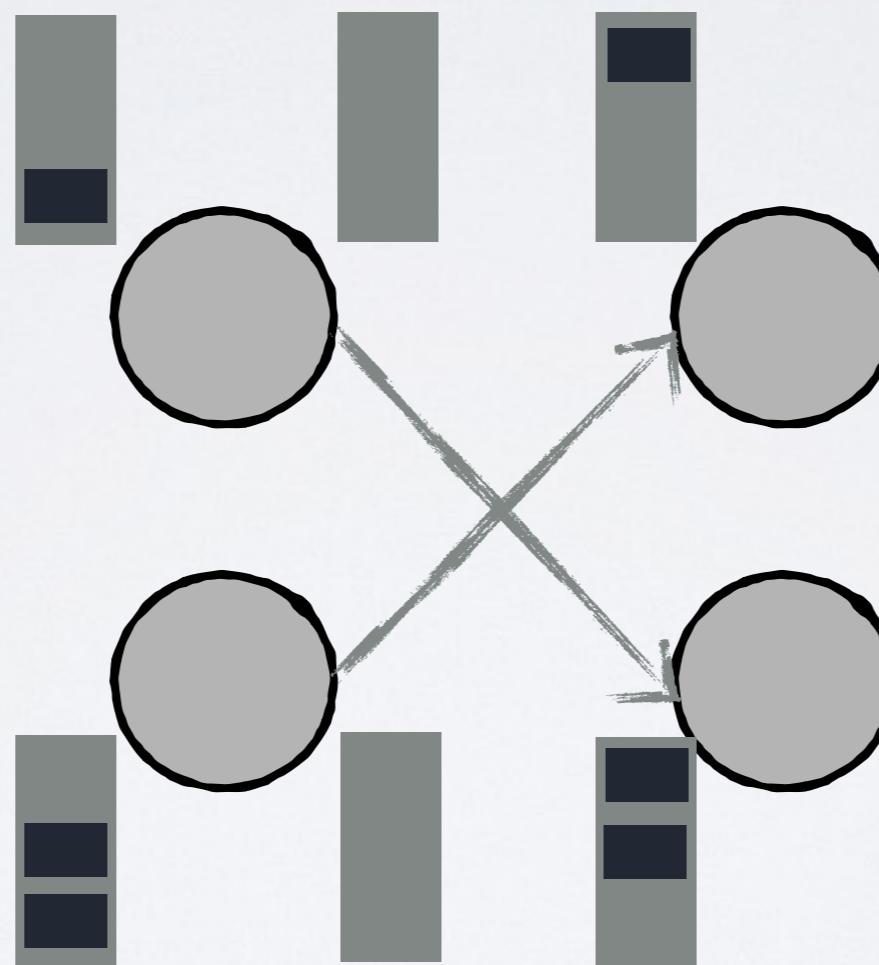
# STREAM ENGINE



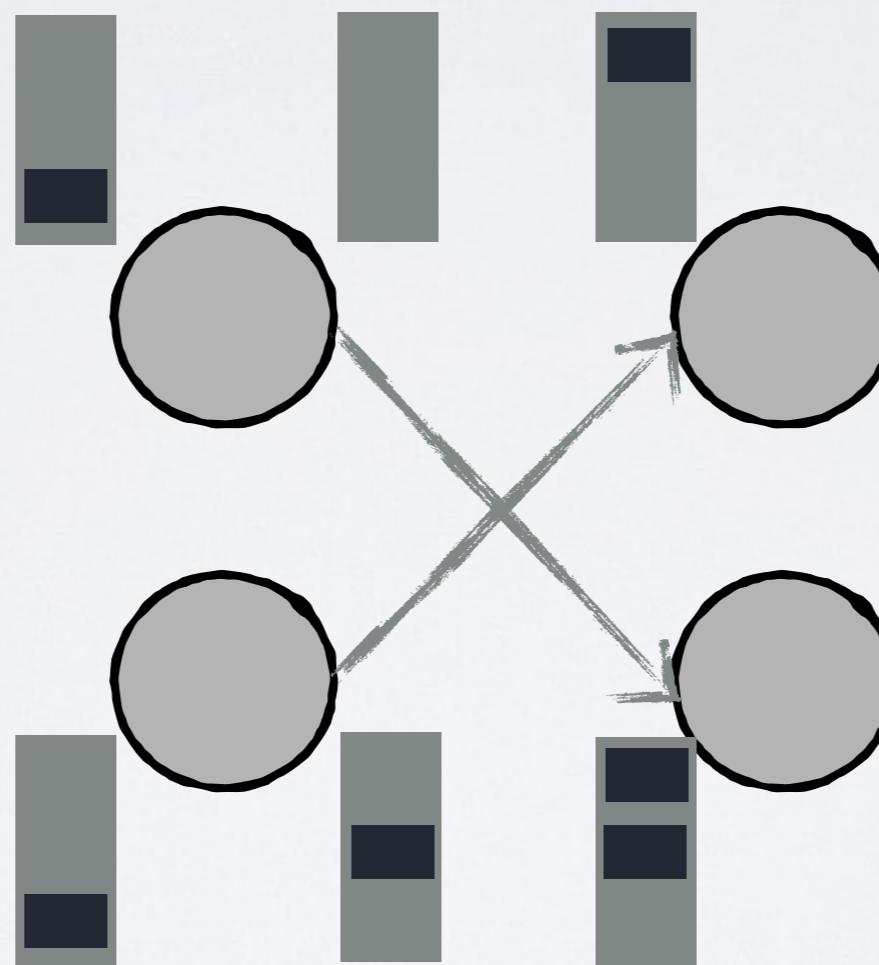
# STREAM ENGINE



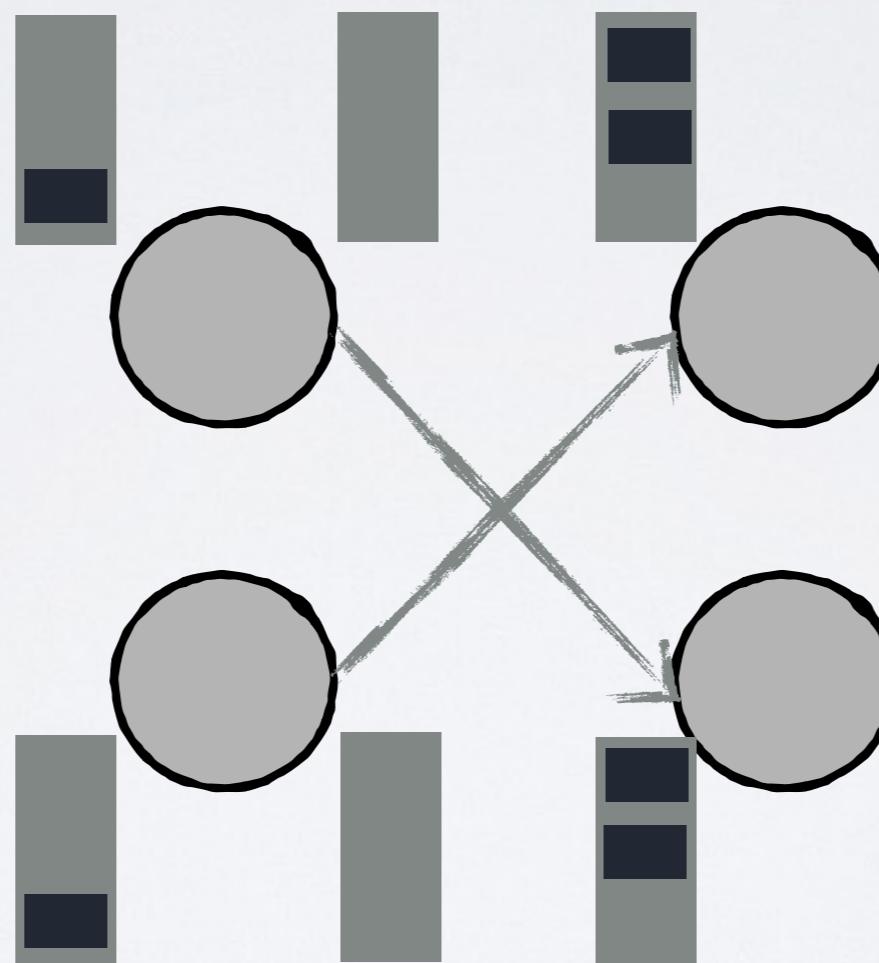
# STREAM ENGINE



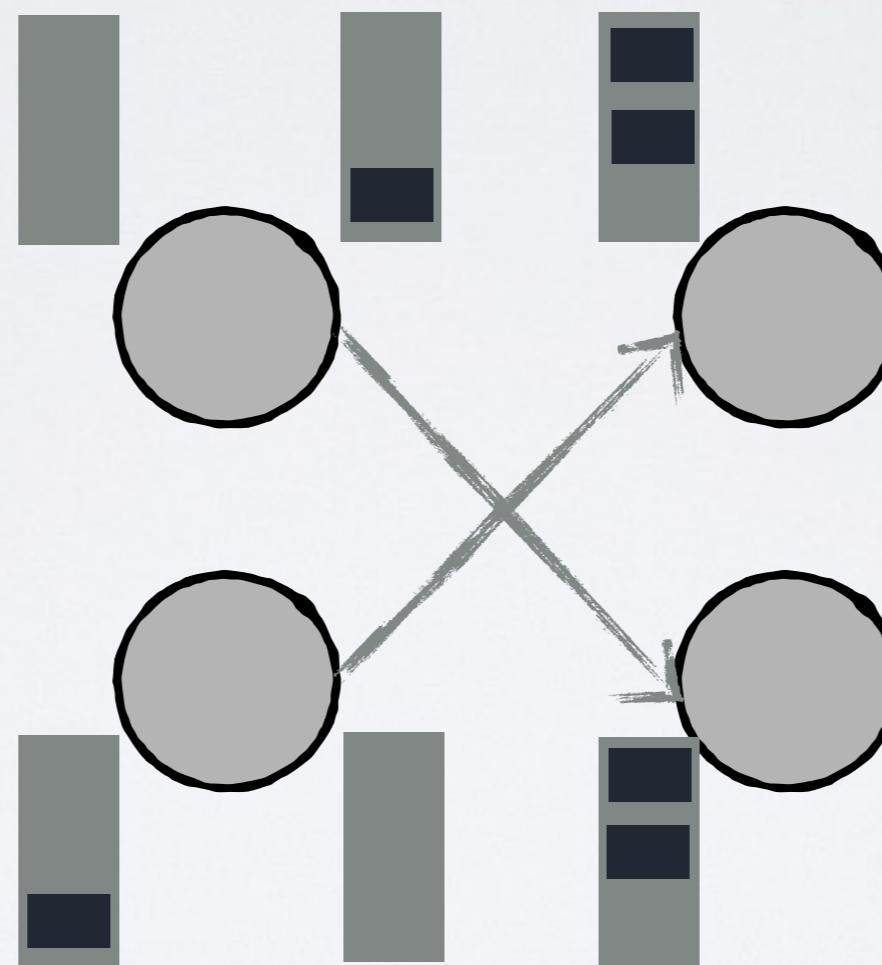
# STREAM ENGINE



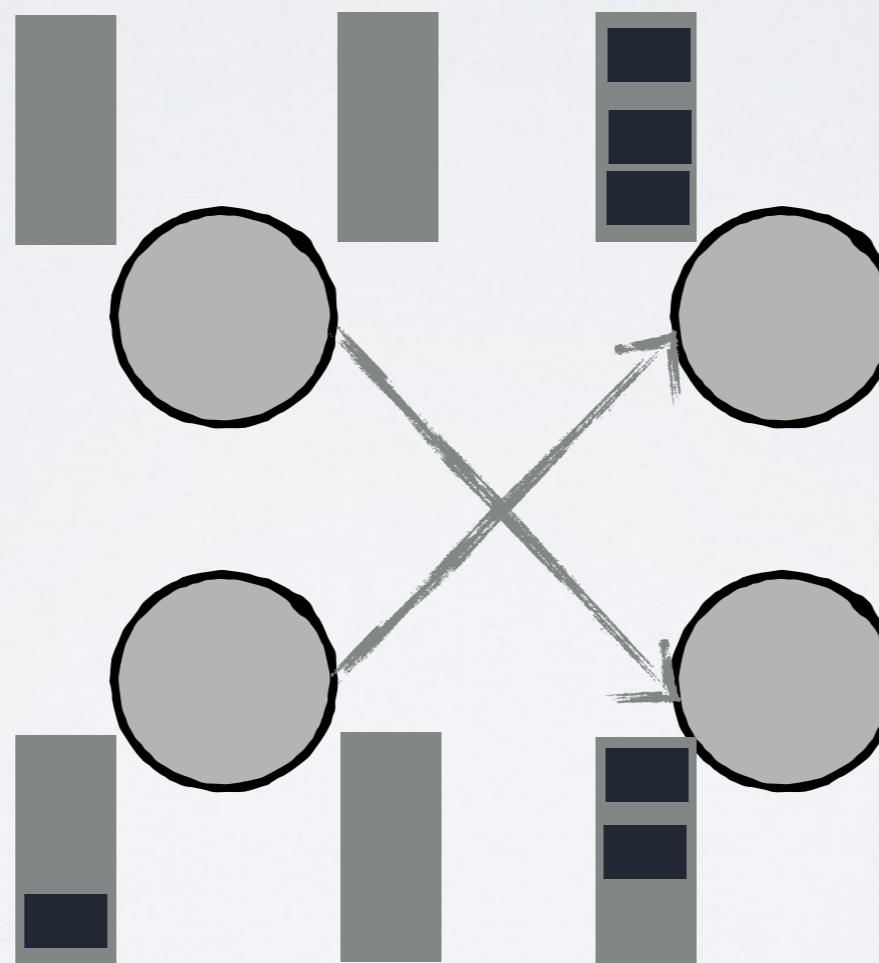
# STREAM ENGINE



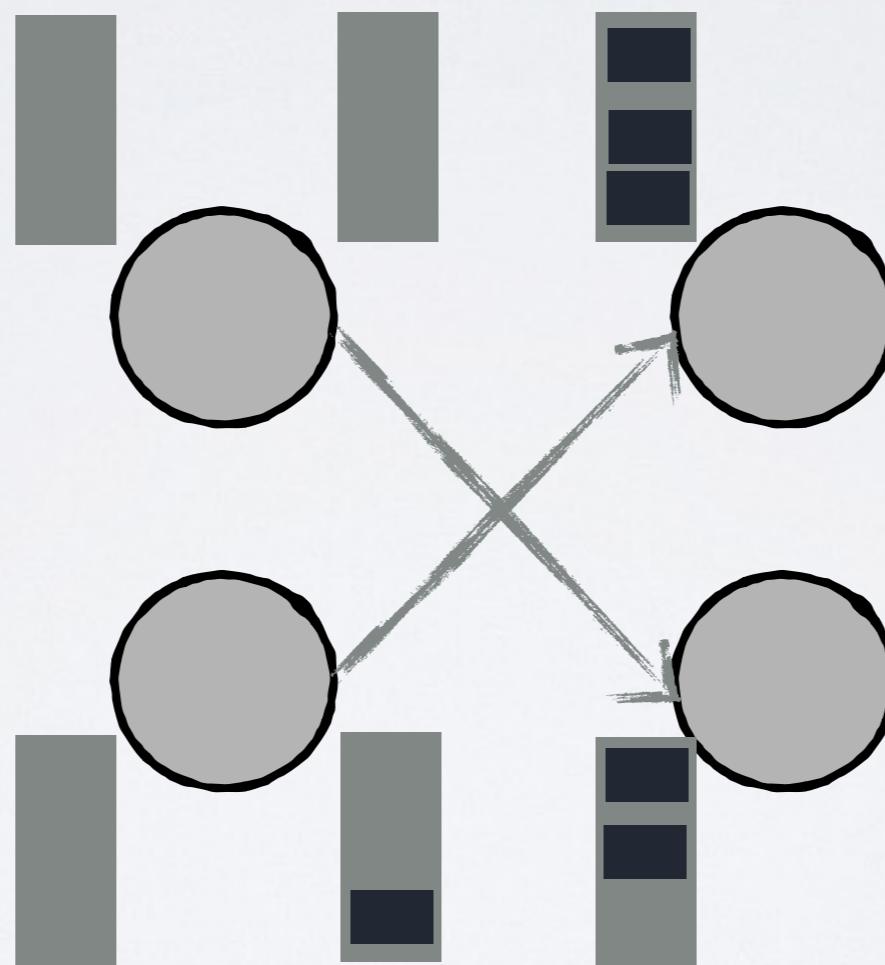
# STREAM ENGINE



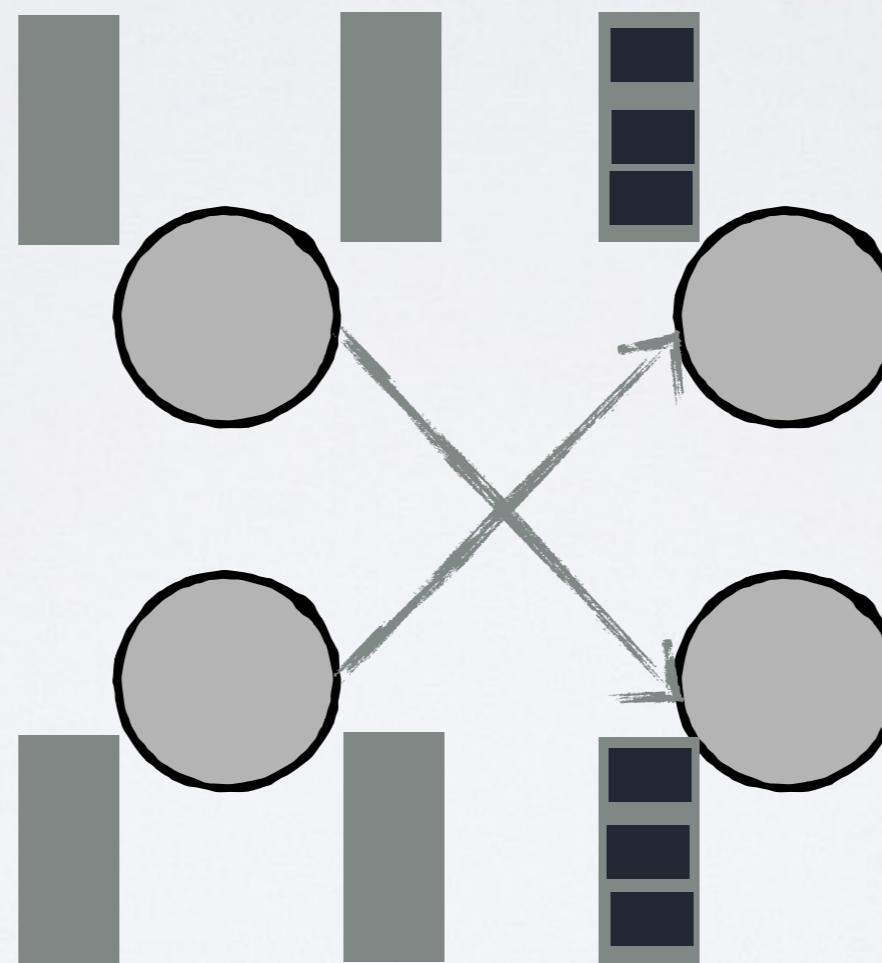
# STREAM ENGINE



# STREAM ENGINE



# STREAM ENGINE



BATCH ENGINE FOR BOUNDED  
DATA AND STREAM ENGINE  
FOR UNBOUNDED DATA, WHY?

# BATCH ENGINE FOR UNBOUNDED DATA

Divide unbounded data into  
bounded data by window.



# STREAM ENGINE FOR BOUNDED DATA

Better resource utility with the  
cost of task level fault tolerant.



# Flink

DataSet(Boun  
ded data)

DataStream(Un  
bounded Data)

Stream Engine

# Spark

StreamRDD(Unbounded Data)

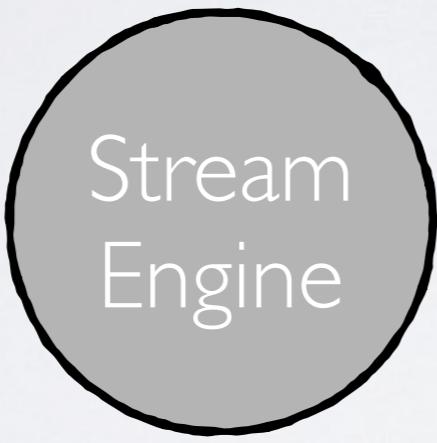
RDD(Bounded data)

Batch Engine

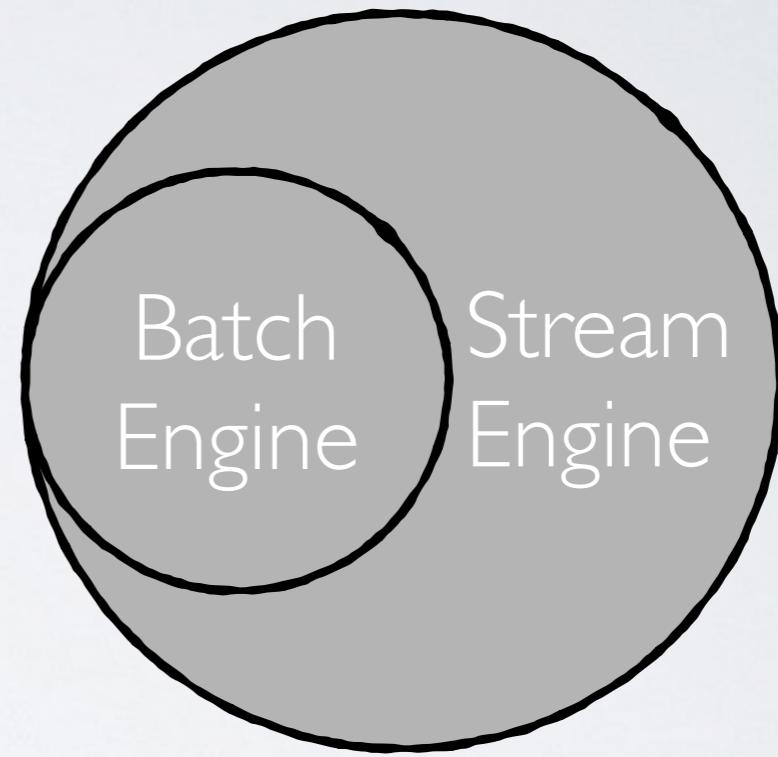
Which one is better?

# FLINK STREAM ENGINE

- Flink Runtime Engine support both Batch and Pipeline mode.
- For DataSet API, use Batch or Pipeline.
- For DataStream API, use Pipeline.



or



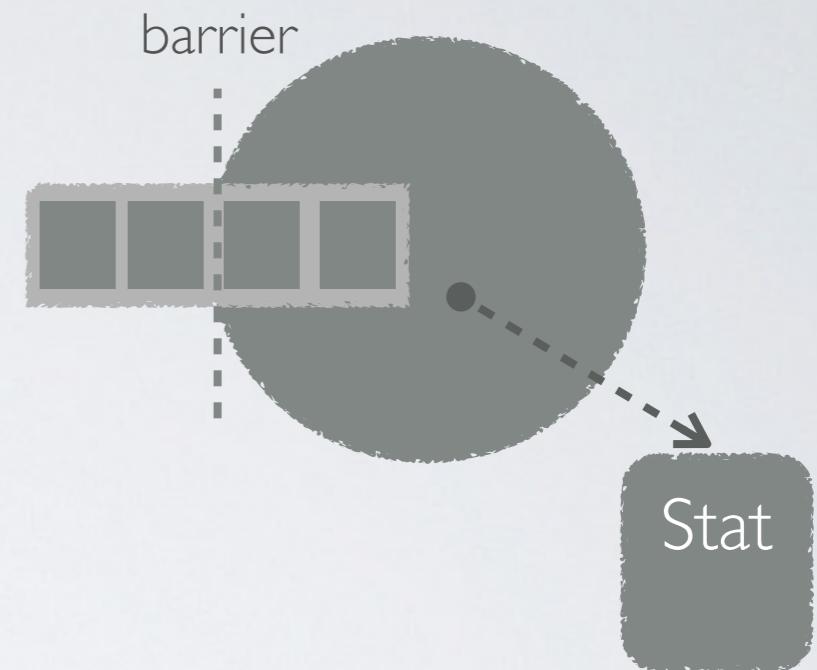
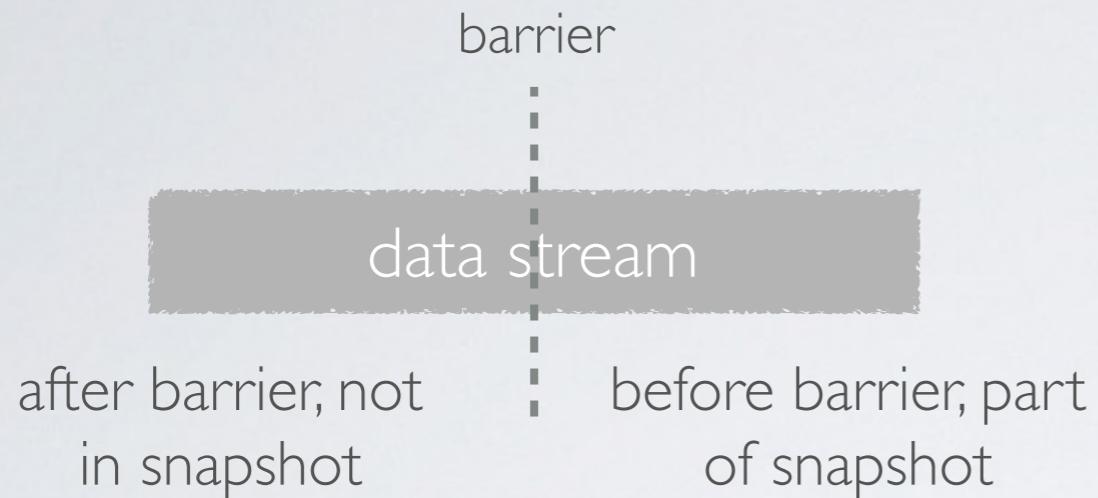
WHAT SHOULD YOU CARE  
ABOUT STREAM?

CORRECTNESS IS THE STAPLE  
FOOD INSTEAD OF DESSERT

Batch Engine intermediate state = all data +  
partial computation

Stream Engine intermediate state = partial data  
+ all computation

- Distributed Snapshot



- Exactly-Once support.
  1. do distributed snapshot.
  2. rewind messages after barrier from source.
  3. support rewinding messages at source side.
  4. support UPDATE at sink side.

WHEN IT'S HAPPENED  
INSTEAD OF WHEN IT'S  
PROCESSED

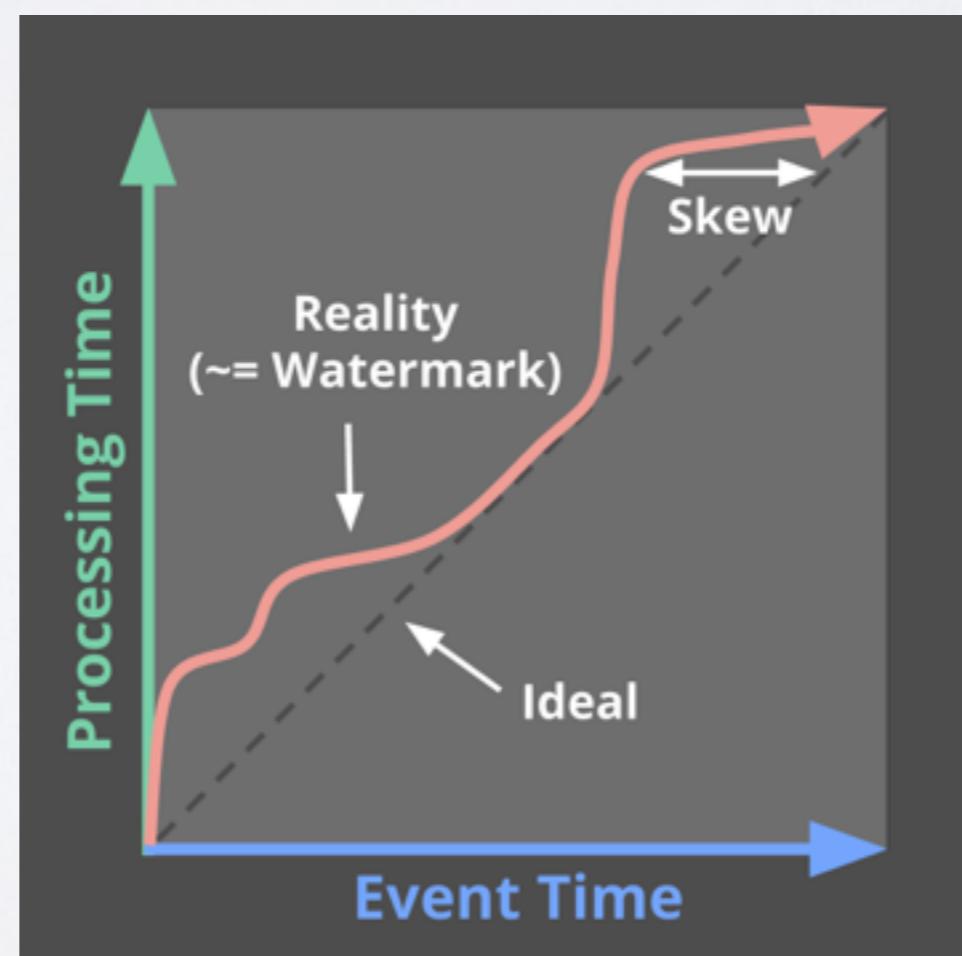
- Process Time

The time at which events are observed in the system.

- Event Time

The time at which events actually occurred.

- Watermark



- Window
  - Process Time/Event Time/Session
  - Tumbling/Sliding
- Trigger
  - Process Time/Event Time/Count
  - Continuous/Purging/Delta
- Evictor
  - Time/Count/Delta

# CEP(COMPLEX EVENT PROCESSING)

```
Pattern<MonitoringEvent, ?> warningPattern =  
Pattern.<MonitoringEvent>begin("First Event")  
.subtype(TemperatureEvent.class)  
.where(evt -> evt.getTemperature() >= TEMPERATURE_THRESHOLD)  
.next("Second Event")  
.subtype(TemperatureEvent.class)  
.where(evt -> evt.getTemperature() >= TEMPERATURE_THRESHOLD)  
.within(Time.seconds(10));
```

# FLINK STREAM SQL

```
val env = StreamExecutionEnvironment.getExecutionEnvironment()

// create some stream in the Streaming API

val stream : DataStream[(String, Double, Int)] = env

.addSource(new FlinkKafkaConsumer(...))

.map { line => parse(line) }

// create Table Environment and register stream with column names

val tabEnv = new TableEnvironment(env)

tabEnv.registerDataStream(stream, "myStreamTab", ("ID", "MEASURE", "COUNT"))

val rTable:Table = tabEnv.sql(

    "SELECT ID, MEASURE FROM myStreamTab WHERE COUNT > 17")
```

THANKS

# REFERENCE

1. <https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>
2. <http://radar.oreilly.com/2014/07/why-local-state-is-a-fundamental-primitive-in-stream-processing.html>
3. <http://flink.apache.org/>
4. <http://spark.apache.org/>